

Churn Prediction With Graphical Models

The goal is to predict if a currently active user will no longer be active next week, within mobile games.

Twitter: @GrimmScientist
Email1: TheGrimmScientist@gmail.com
Email2: allen@playhaven.com
Phone: 724-312-4157

Presenter: Allen Grimm
From: PlayHaven
On: Nov 9, 2013
At: PyData, NYC

Structure of Presentation

- Overview of PlayHaven
- Definition of Churn
- Our Graphical Model Framework
- Experiment 1 – Predictions for one game.
- Experiment 2 – Predictions across the full network.

PlayHaven:

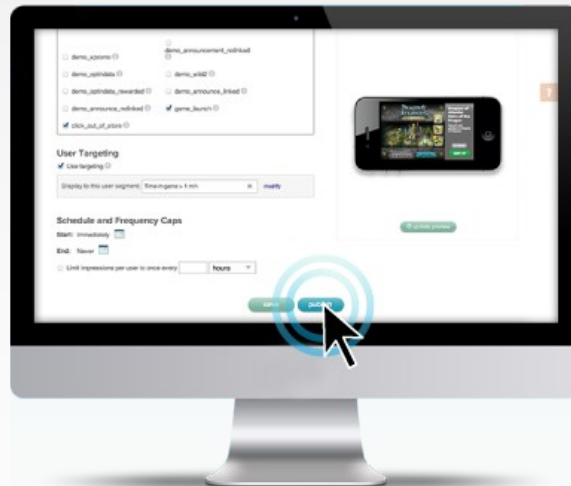
- What we do
- How data science can help

Introducing: PlayHaven

The **Business Engine** to help mobile game developers acquire, engage, and monetize players.



SDK



Dashboard



Action

Opportunities for Data Science at PH

- Automate everything
- Response modeling
- Predictive segment dimensions
- Advanced analytics

Player Churn:

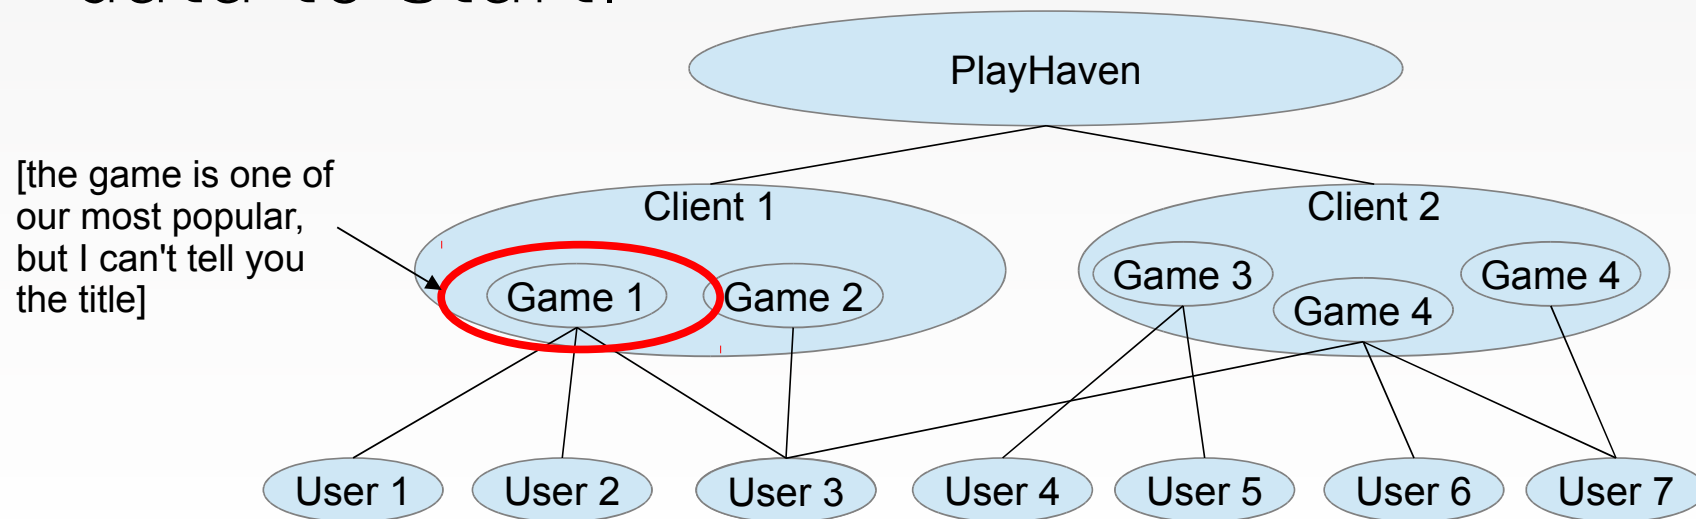
- Let's make a precise definition

Introducing: Player Churn

- There is a precise cost of acquisition for each user (it's never free).
- Can we detect when a user is about to leave?
 - (if so, we can presumably do something about it)
- Churn: When a device used to be active and is no longer active.

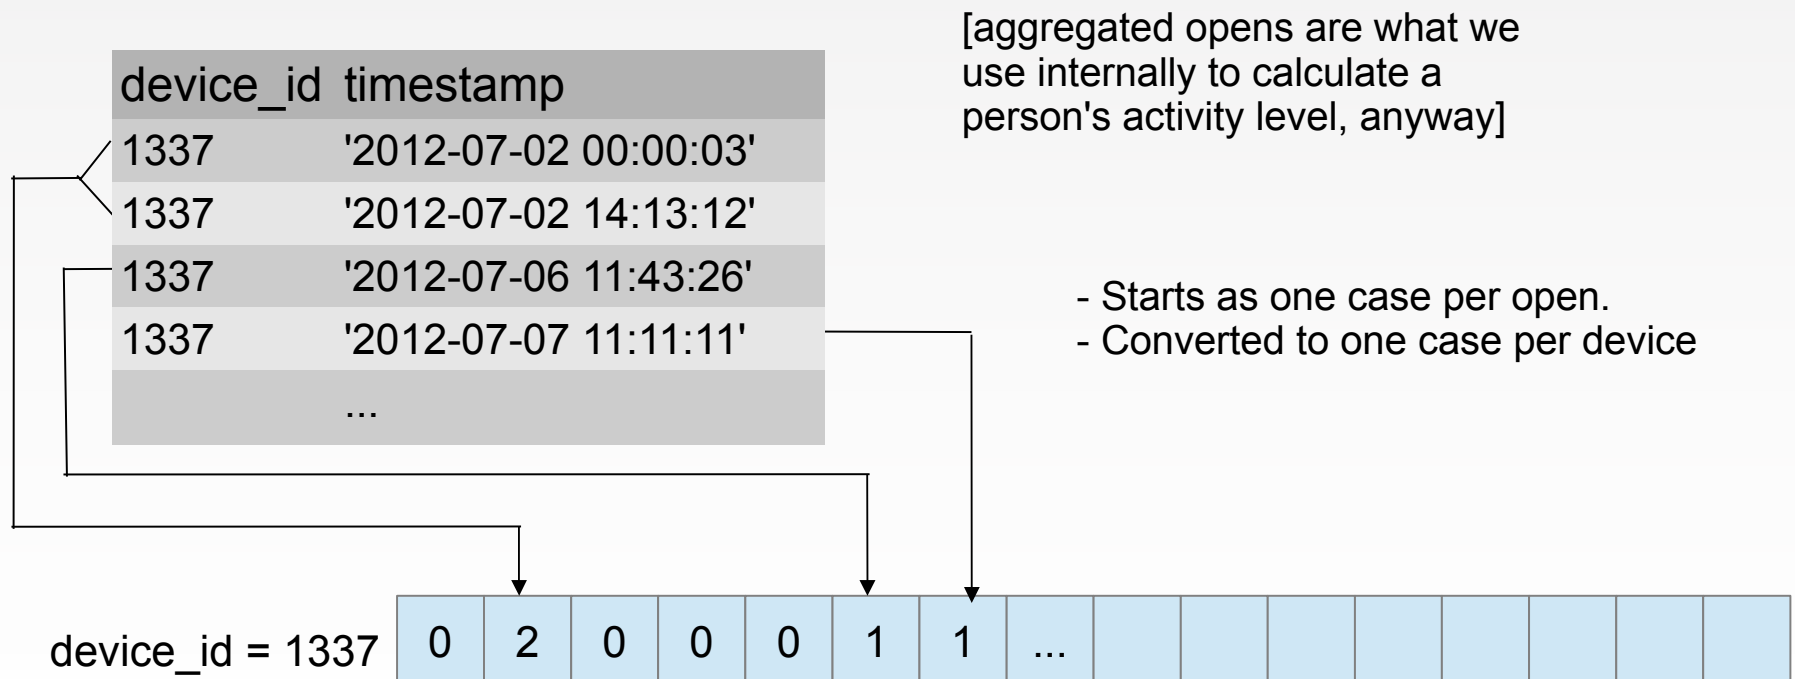
Let's start with just one game

- We have lots of publishers with lots of games.
- Let's just pick one game's worth of data to start.



Our Data: Aggregated Game Opens

- We have lots of event types, but let's start simple.





Churn – a user who is no longer active



Will Churn

Will Churn Soon



Any
activity



Medium
activity



High
activity





Churn – a user who is no longer active



Just Churned

Will Churn Soon



Any activity



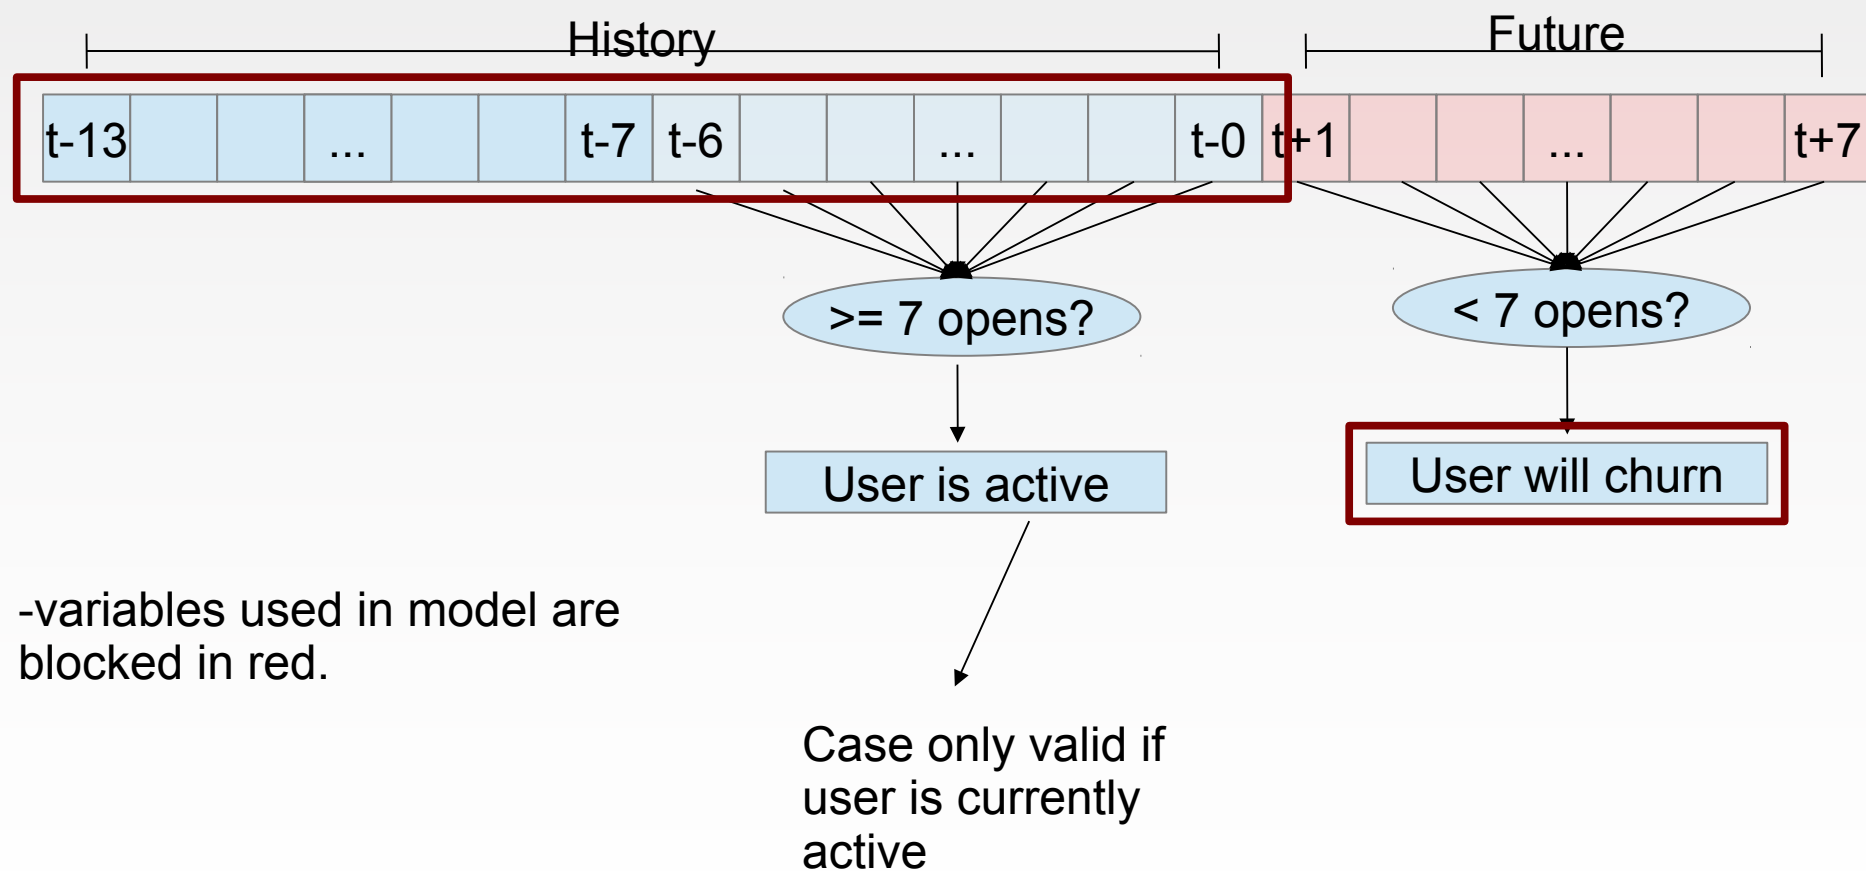
Medium-active



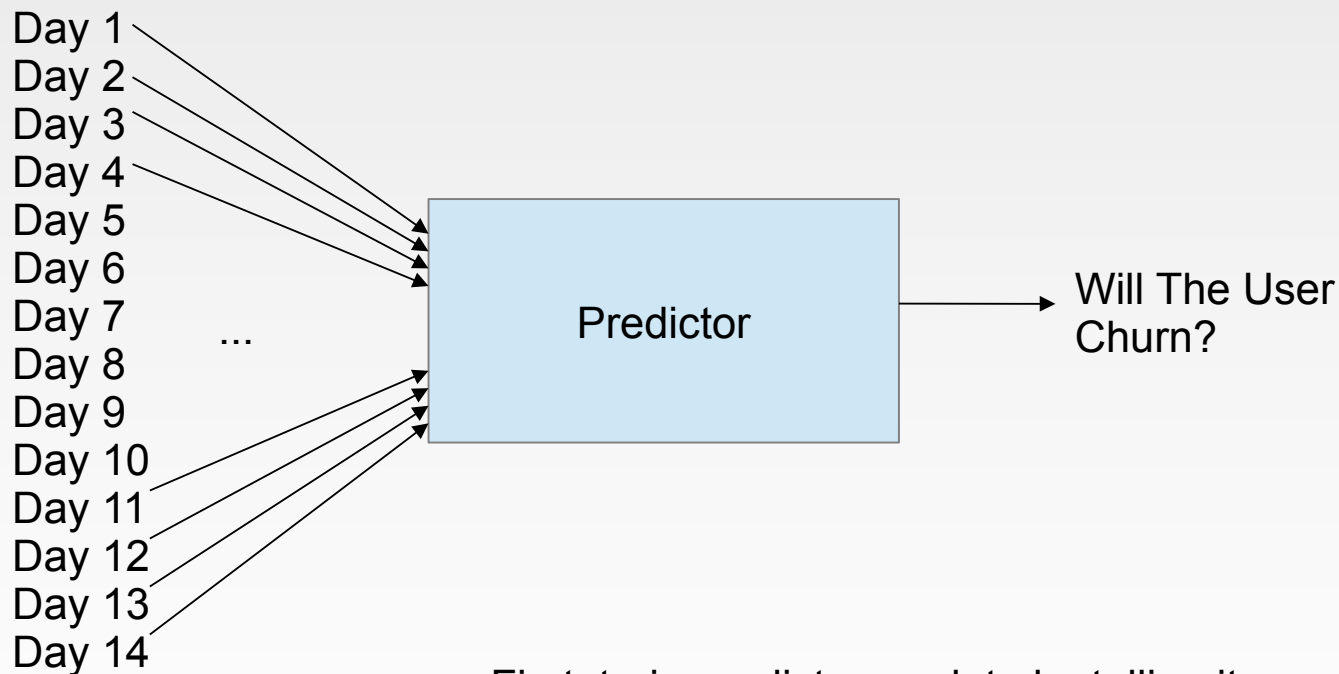
Extremely active



Churn – A Single Sample of Data



Our Black Box for Supervised Learning



- First, train predictor on data by telling it both the input and output.
- Second, test the predictor by only giving it inputs and checking what output it gives.

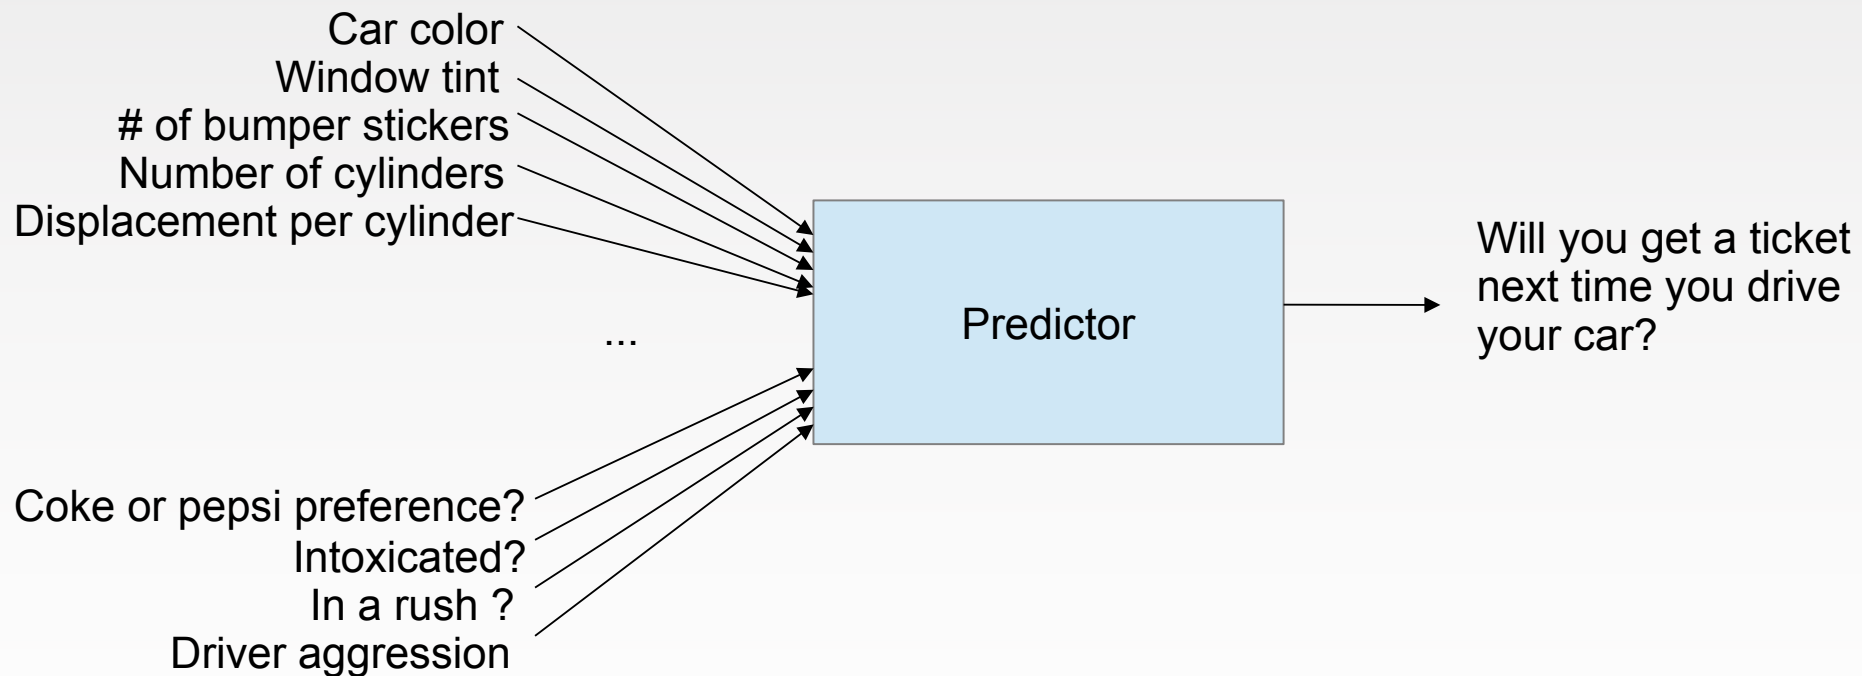
Our Graphical Model Framework:

- Reconstructability Analysis
- Representationally: sets of sets of statistically connected variables

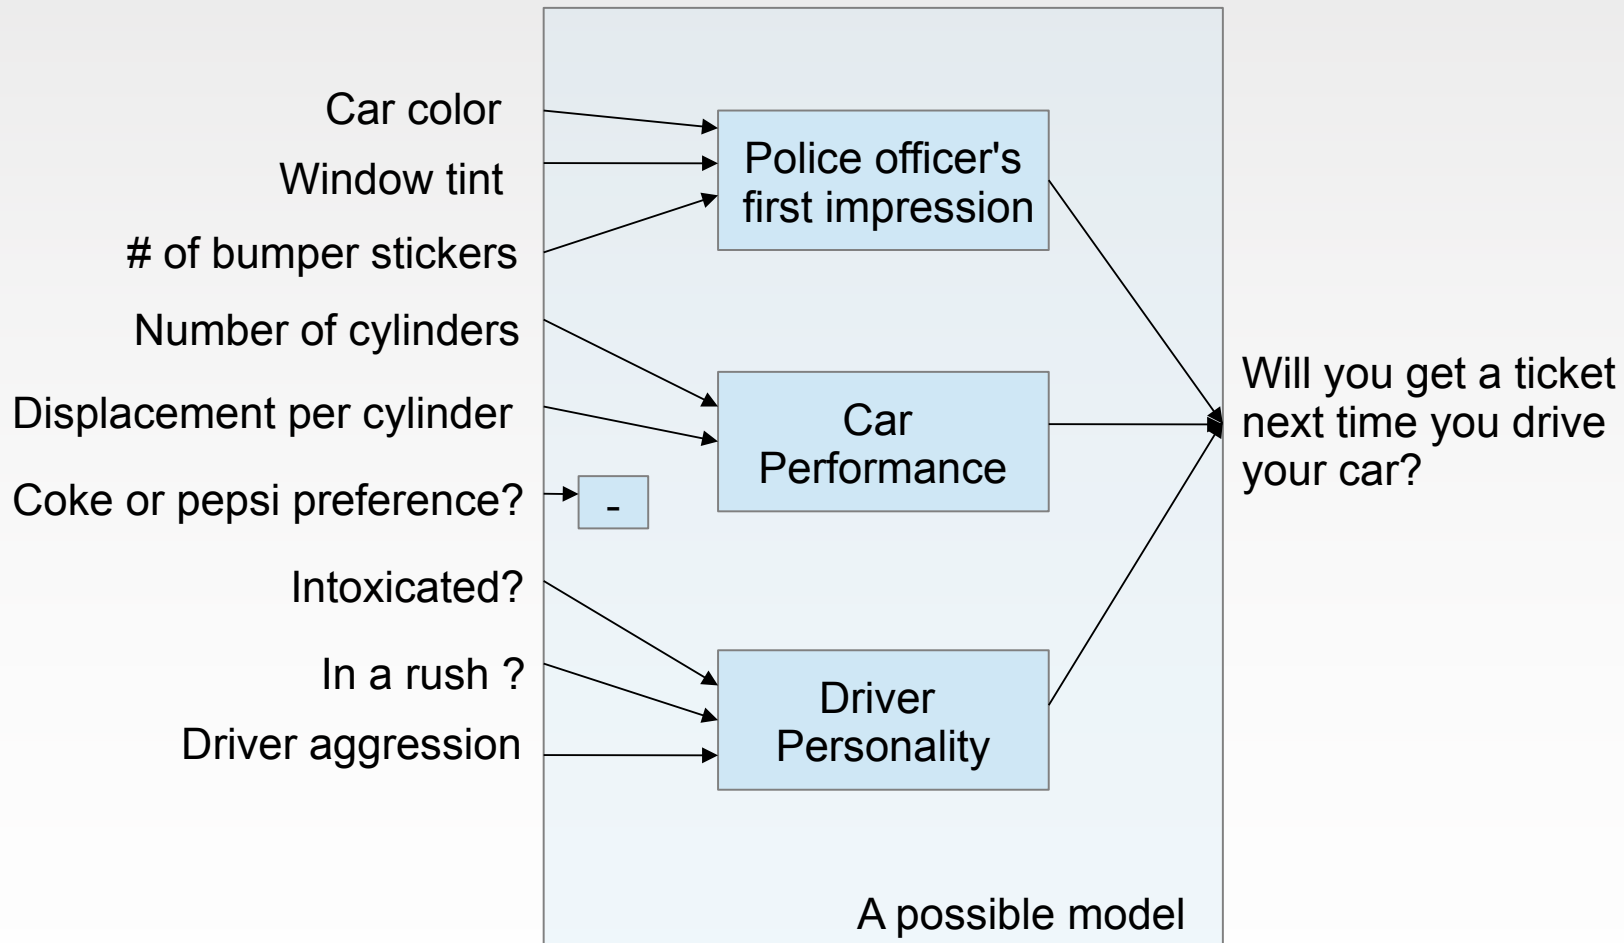
Reconstructability Analysis

- A graphical model framework with heavy overlap with Bayesian Networks and Loglinear Models.
- Uses **contingency tables to build conditional probabilities.**
- An RA model is a set of components
- Each component is a set of fully connected variables

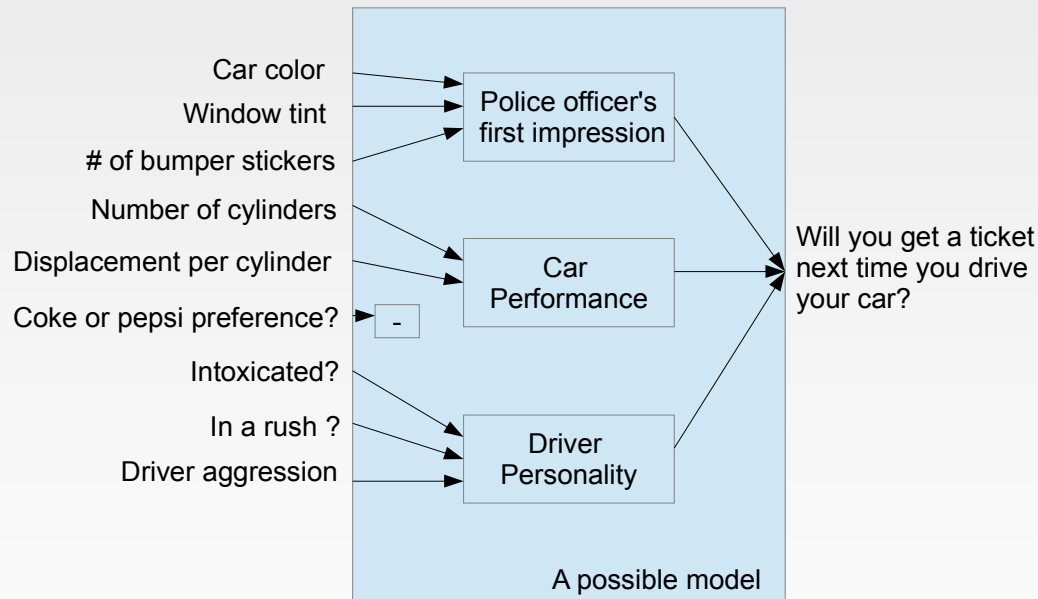
Let's Start With An Example...



Example: Underlying Structure



RA – Representation Example



Long Variable Name	Letter
Car Color	A
Window Tint	B
# of Bumper Stickers	C
Number of Cylinders	D
Cylinder Displacement	E
Coke or Pepsi?	F
Intoxicated?	G
In a Rush?	H
Driver Aggression	I
Will You Get A Ticket?	Z

Component Structure:

“First Impression” → K1 → ABCZ

“Car Performance” → K2 → DEZ

“Driver Personality” → K3 → GHIZ

Model Structure:

Model → K1:K2:K3

→ ABCZ:DEZ:GHIZ

RA – “Goodness” example

Example dataset:

	F = Coke		F = Pepsi	
	A = red	A = blue	A = red	A = blue
Z=yes	98	4	102	6
Z=no	8	90	4	88

Variable Name	Letter
Car Color	A
Coke or Pepsi?	F
Ticket?	Z

[Projections]

“AZ” Projection:

	A = red	A = blue
Z=yes	200	10
Z=no	12	178

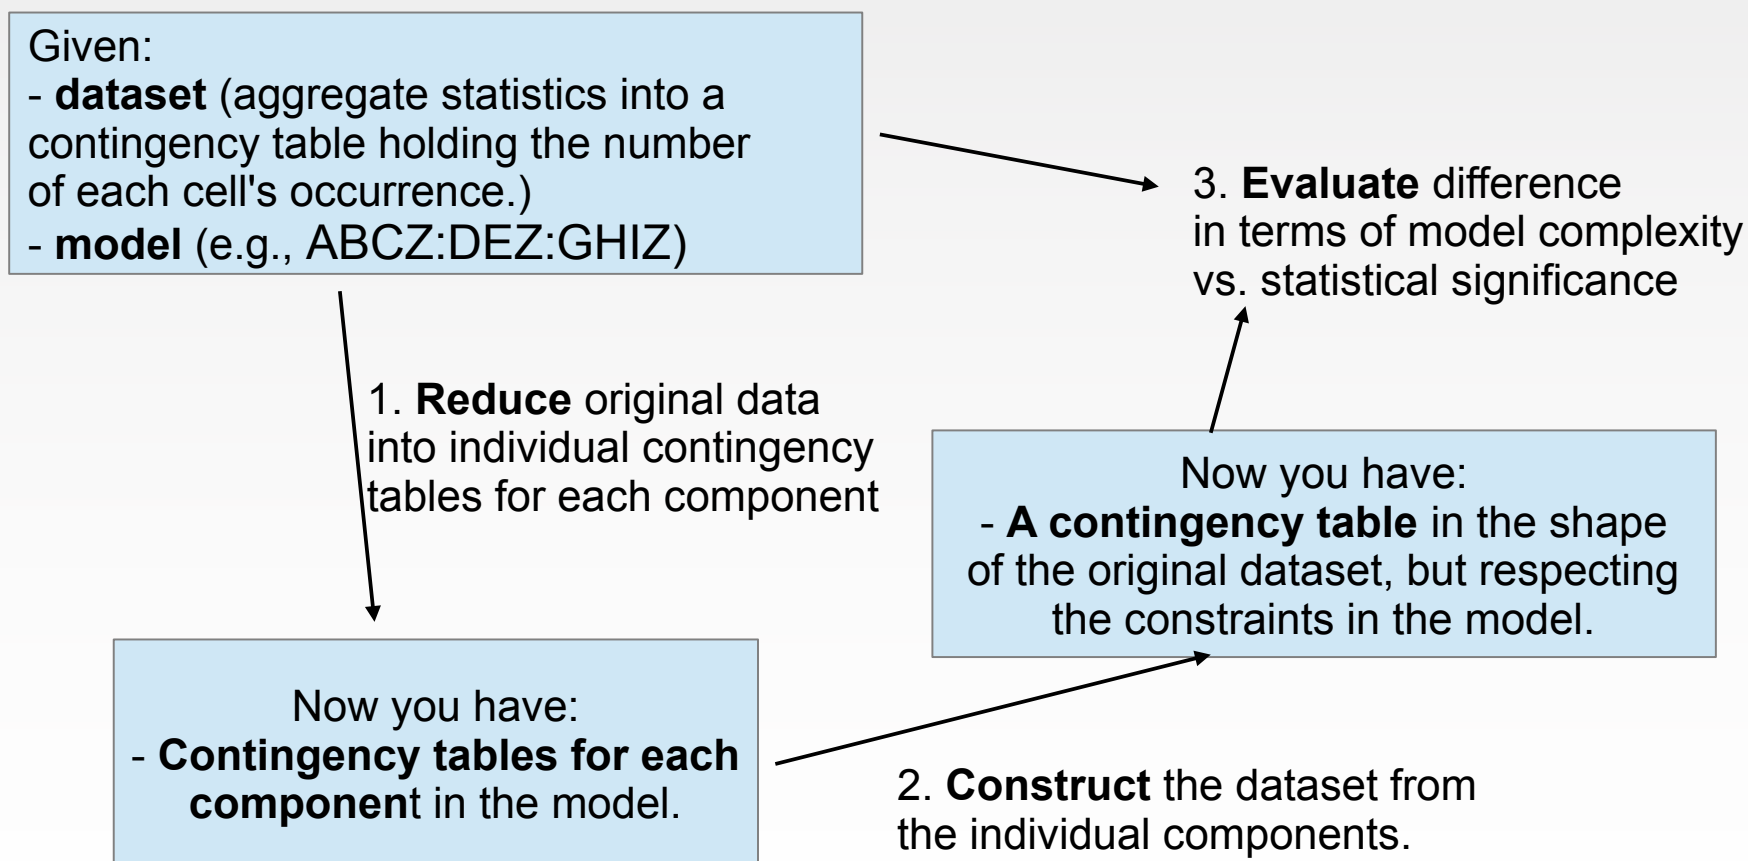
“better”, since it contains more information

“FZ” Projection:


	F = Coke	F = Pepsi
Z=yes	102	108
Z=no	98	92

“worse”, since it loses most of the information

RA – Model Evaluation Process




RA – Lattice of Structures



Degrees of freedom (dF) increases as you climb the lattice

Likelihood (L^2) increases as you descend the lattice



Ways to balance dF and L2:

- Information Criterion:
 - **Bayesian** (BIC)
 - Akaike
 - Deviance
- Incremental alpha
- % correct on test set

Experiment 1:

- Search for a model
- Use that model as predictor in a game

Introducing: Occam

- Occam is a web interfaced tool built and maintained at Portland State University for running RA searches and evaluating RA models on arbitrary datasets.
- General information at <http://dmm.sysc.pdx.edu>
- Tool at <http://dmm.sysc.pdx.edu/weboccam.cgi>
- Note: we are currently discussing “variable based” RA.

Search Results

“Loopless” Model:

Level	model	dDF	%dH(DV)	dBIC	%c (train)	%c (test)	%cover (train)	%cover (test)
4	IV:KLMNO	624	21.2287	271665.94	78.77	90.5	100	100

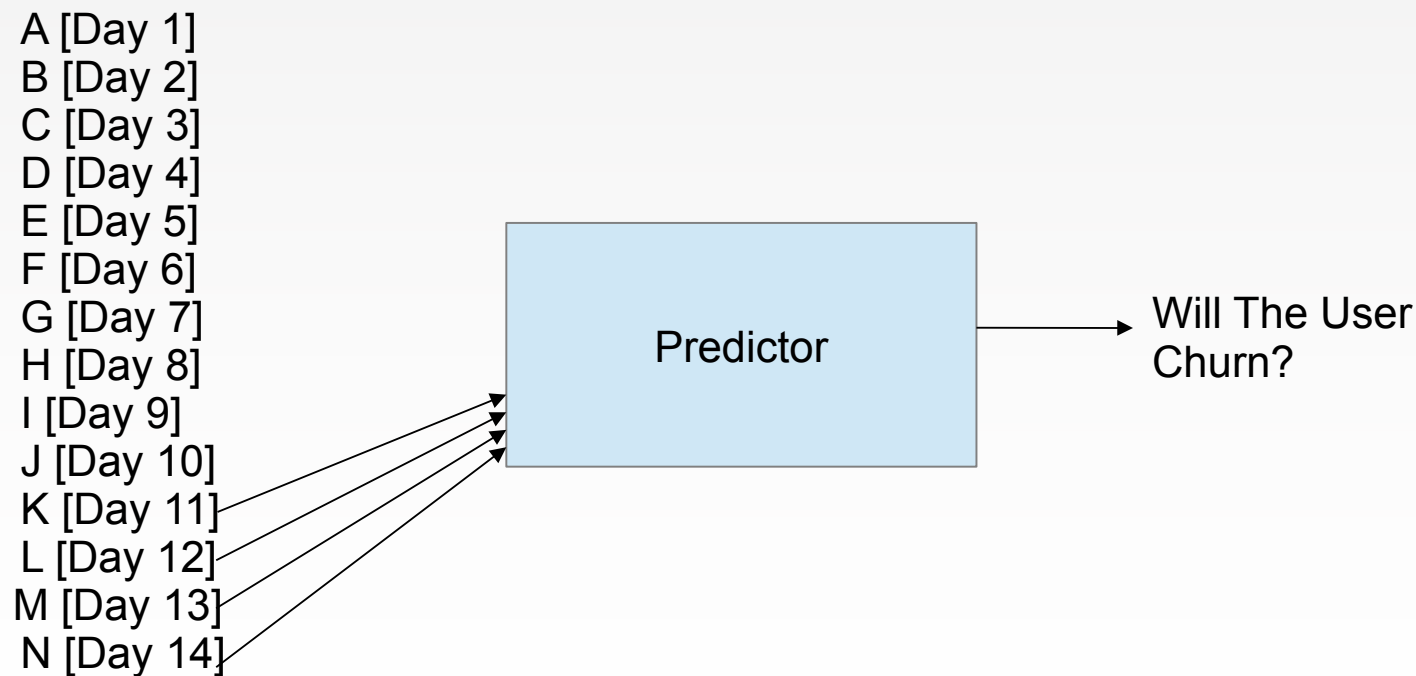
“Loopy” Model:

Level	model	dDF	%dH(DV)	dBIC	%c (train)	%c (test)	%cover (train)	%cover (test)
20	IV:BIO:BKO:CHO:CJO:HIO:ILO :JKO:KMO:LMO:LNO:MNO	212	23.5297	307800.93	79.50	87.78	79.50	76.56



Our Predictor: KLMNO

- Loopless models run way faster than loopy models



Results: Confusion Matrix

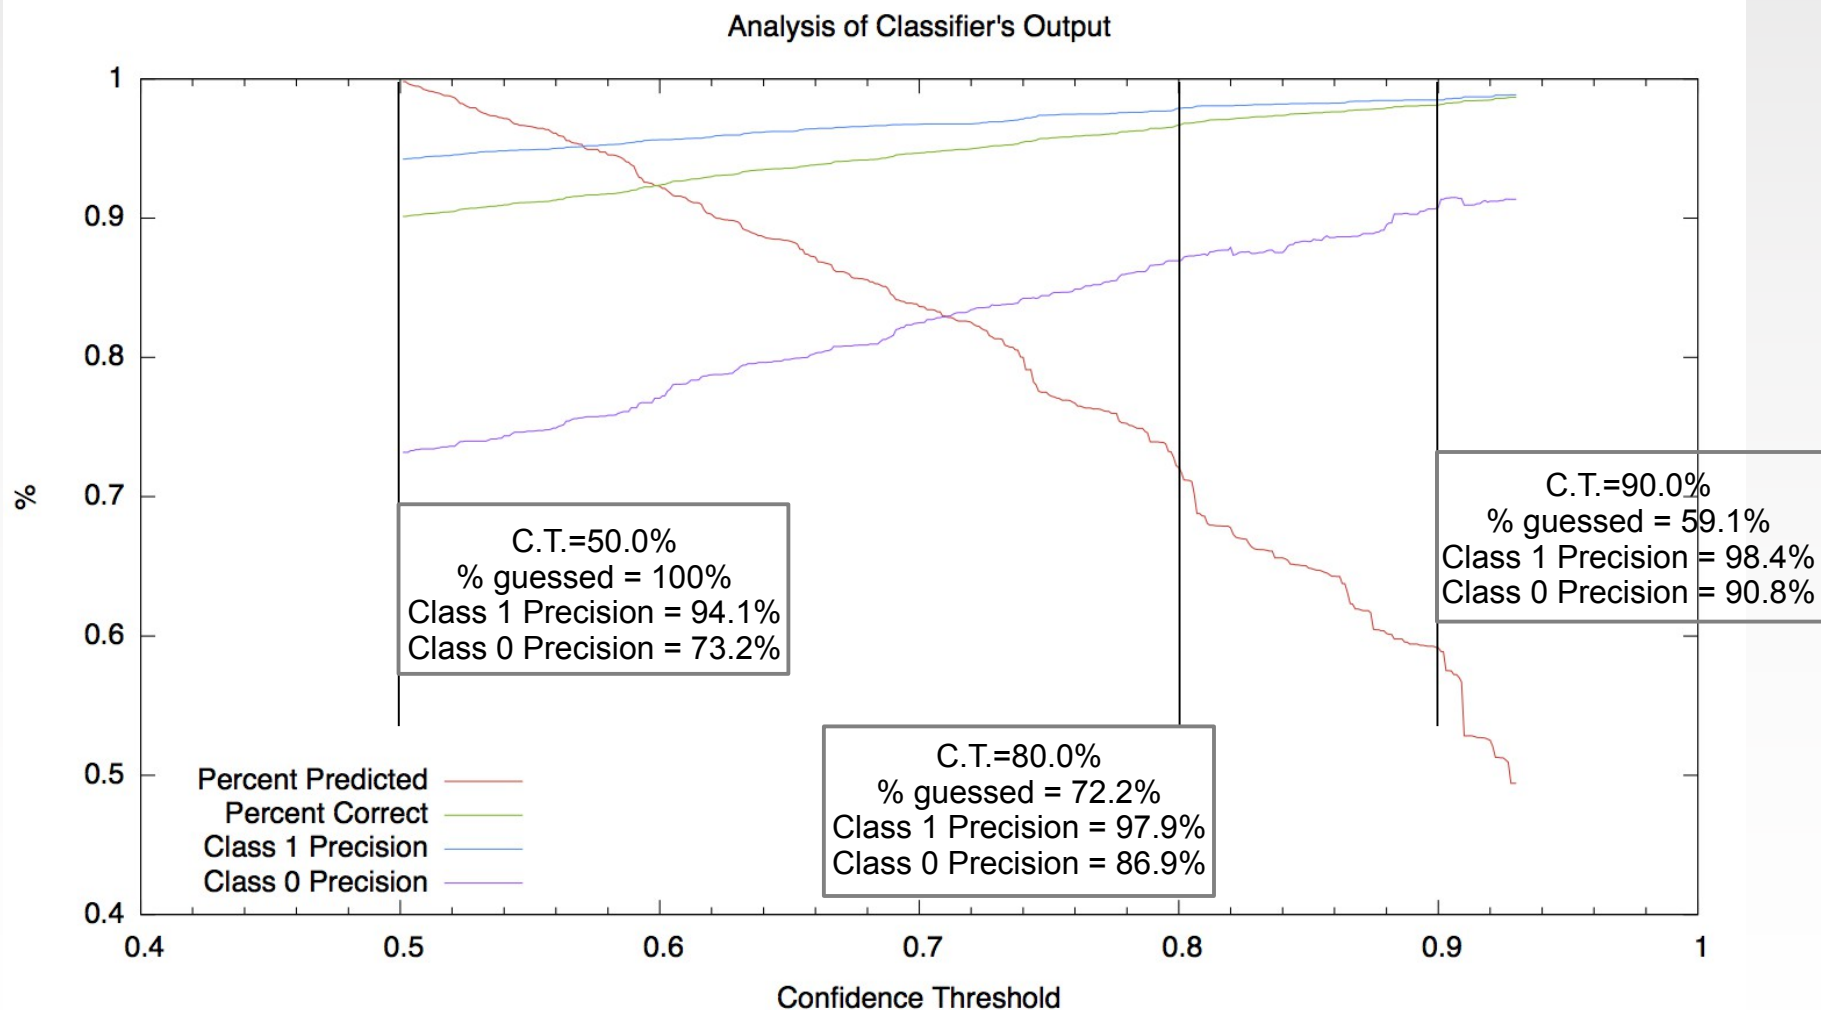
- Recall the best-by-BIC, loopless, directed model: “KLMNO”
- Train/test sets generated by picking two timesteps a month apart.

	Predicted to Churn	Predicted to Not Churn
Actual Will Churn	818,274	56,481
Actual Will Not Churn	50,970	154,183

Works out to be:

- 90.04% prediction rate,
- 93.54% will-churn precision
- 75.16% will-not-churn precision

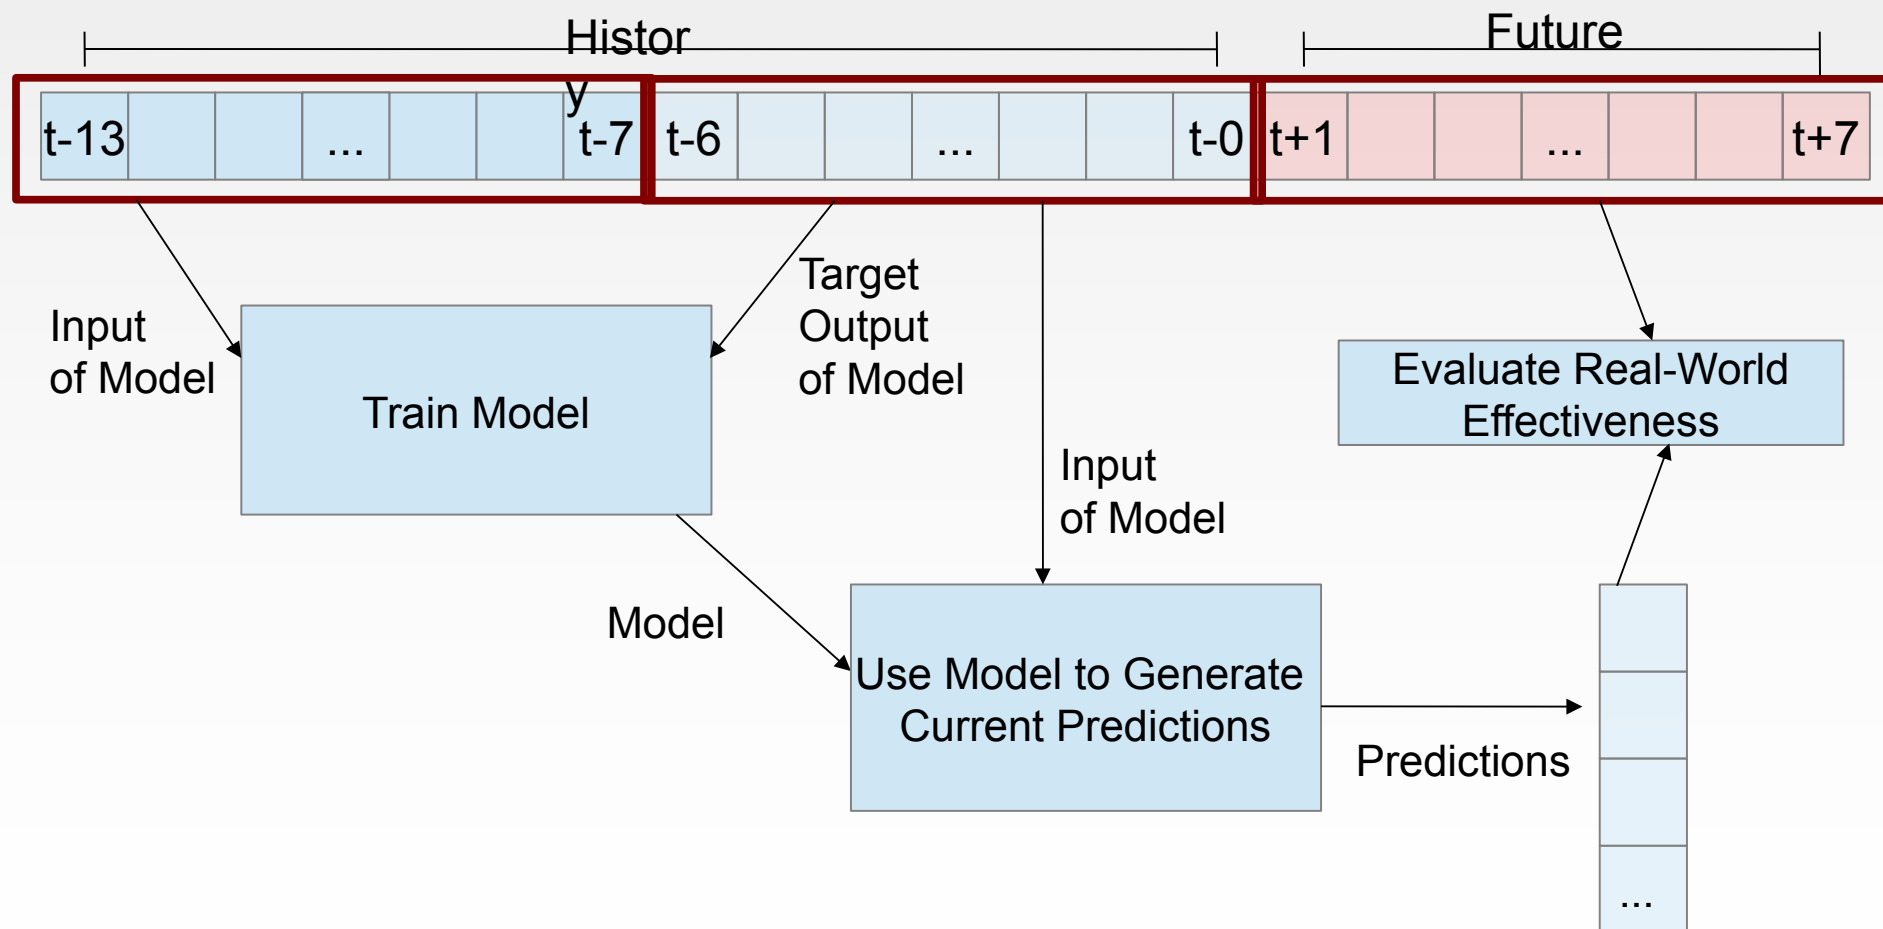
The Confidence Threshold



Experiment 2:


- Discuss productized implementation
- Evaluate performance on full network

Rolling Window For Model Training



Compute Time On Initialization

"Raw" data files




- (device,game,count) tuples.
- Exported from Redshift DW.
- 5 minutes per day of data.
- 11.5MM, 200Mb each

User History




- aggregated history for each user.
- 35 – 70 minutes per day aggregated.
- 78MM rows, 7Gb

Model



- Full contingency table
- 30 minutes
- 800 rows, 80Kb

Predictions



- (device,game,prediction) tuples
- 30 minutes
- 9.5MM rows, 300Mb

- Total Processing time: around 13 hours on amazon's "m1.large" instance.

Compute Time To Iterate One Day

Grab next file:

dN+1


User History



Model



Predictions



- (device,game,count) tuples.
- 5 minutes for one day of data.
- 11.5MM, 200Mb

Total time: about 2 hours

- backup table: 20 min
- drop last day: 35 minutes
- add new column: 10 minutes
- populate new column: 45 minutes
- 78MM rows, 7Gb
- Full contingency table
- 30 minutes
- 800 rows, 80Kb
- (device,game,prediction) tuples
- 30 minutes
- 9.5MM rows, 300Mb

- Total Processing time: around 3 hours on amazon's "m1.large" instance.

Performance On Full Network

On **2012-10-26**:

8,582,225 active device-game tuples.


Game id not considered in prediction.

	Predicted to Churn	Predicted to Not Churn
Actual Will Churn	2,911,865	1,130,127
Actual Will Not Churn	1,120,336	3,419,897

Works out to be:

- 73.78% prediction rate,
- 72.04% will-churn precision
- 75.32% will-not churn precision

Performance Across Time



Next Steps

- It works well. It's fast. ... and there's still a lot of work to do.
- To do:
 - Compare with other methods.
 - Add variables (day of week, game genre, player type, ...)
 - Evaluate performance over the last year.
 - Re-run search on full-network dataset
 - Optimize code performance