Let's Do Data Science

By: Allen Grimm
www.TheGrimmScientist.com
@GrimmScientist
TheGrimmScientist@gmail.com

Goals:

- Information theory crash course
- Cross Validated data as a test run.
- A few sample models
- Application of a model

A Coin Flip

Balanced

	Heads	Tails
Probability	0.5	0.5

Biased

	Heads	Tails
Probability	0.2	0.8

One Side at a Time

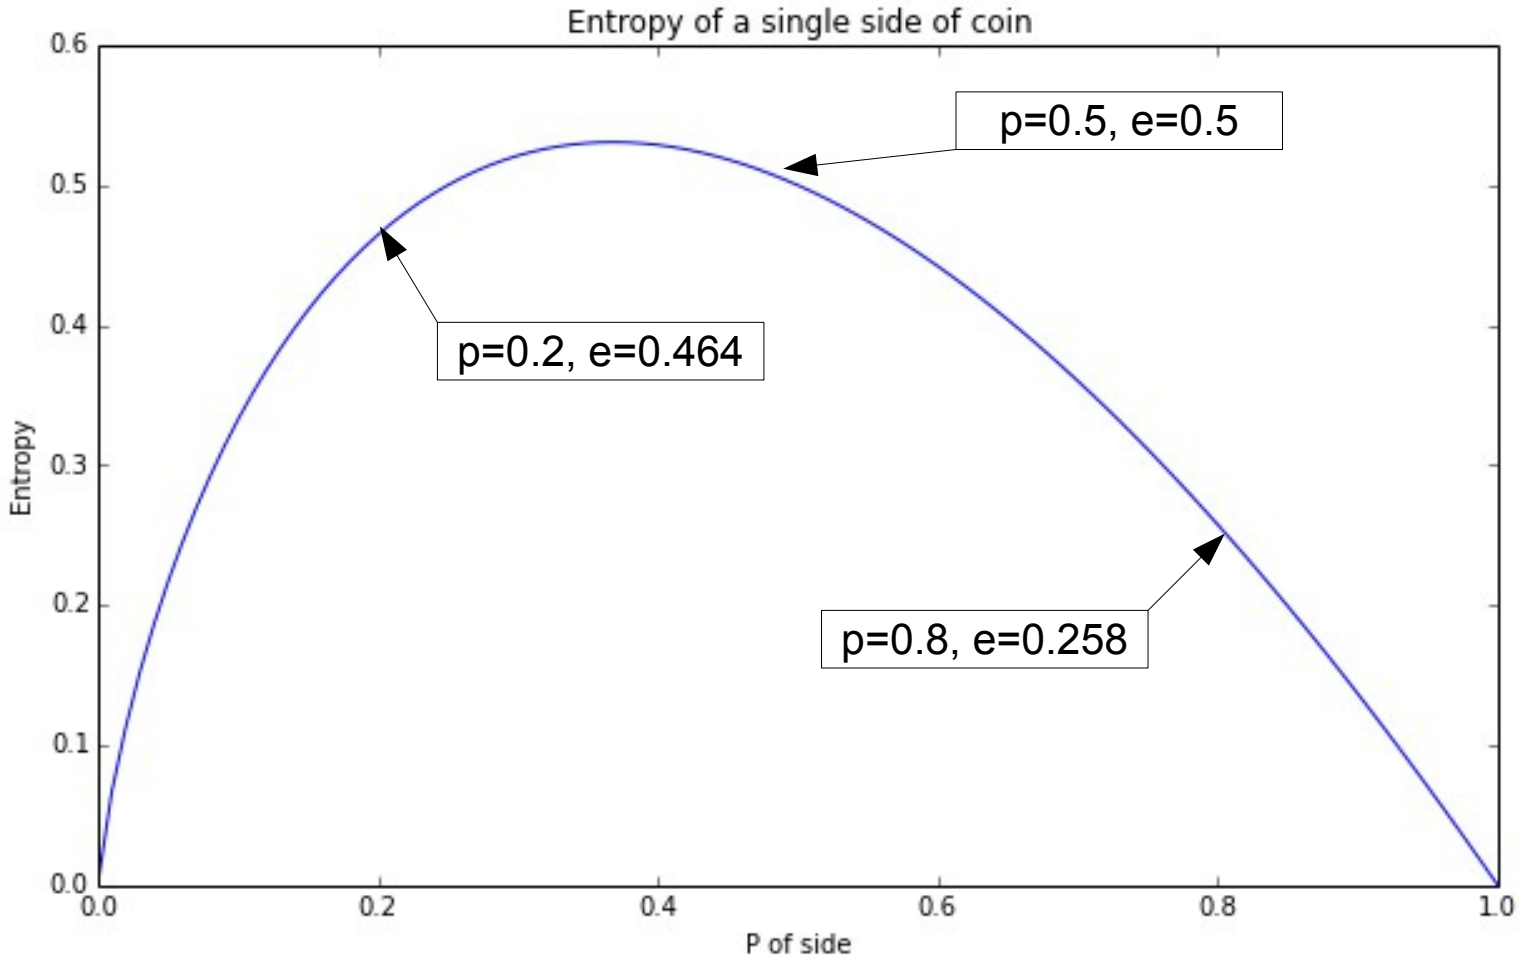

$entropy_i = p_i * \log_2(p_i)$

Both Sides Together

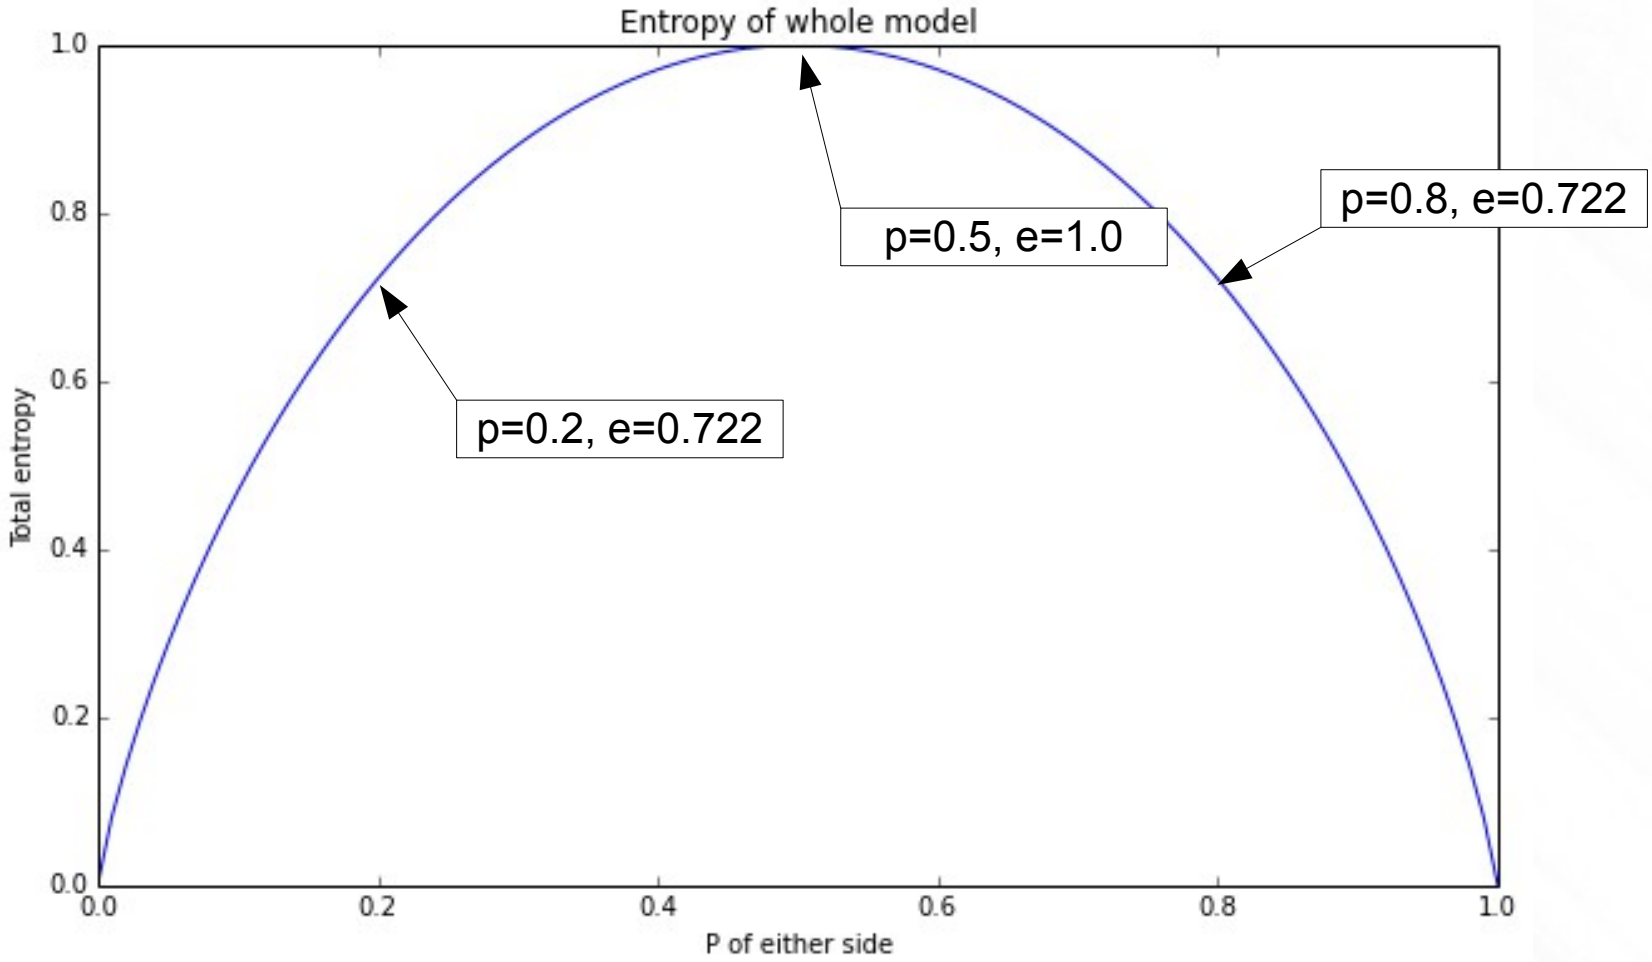

A Coin Flip

Balanced

	Heads	Tails
Probability	0.5	0.5

Total entropy: 1.0

Biased

	Heads	Tails
Probability	0.2	0.8

Total entropy: 0.722

Uniform probabilities always result in maximal entropy.
=
The more biased the probabilities in a model, the more you know about the system ahead of time.

The Data

- Cross Validated post data
- 68,386 posts
- Variables:
  - Score “S”
  - Favorite count “F”
  - Answer count “A”
  - Comment count “C”
  - Body length “B”

Two Comparable Models

Model “SF”

Score	F = 0	F > 0
S < 0	0.0112	0.0005
S = 0	0.1916	0.0086
S > 0	0.6592	0.1289

Favorites

entropy(SF) = 1.37

Model “SA”

	A = 0	A > 0
	0.0062	0.0054
	0.1522	0.0480
	0.5285	0.2596

Answers

entropy(SF) = 1.70

Models of Different Sizes

Model "SF"

Score	F = 0	F > 0
S < 0	0.0112	0.0005
S = 0	0.1916	0.0086
S > 0	0.6592	0.1289

Favorites

entropy(SF) = 1.37
degrees of freedom = 5

Model "SB"

Score	B < 50	50<=B<100	B <= 100
S < 0	0.0045	0.0035	0.0036
S = 0	0.0490	0.0582	0.0930
S > 0	0.1156	0.1910	0.4816

Body Length

entropy(SF) = 2.18
degrees of freedom = 8

Model Selection

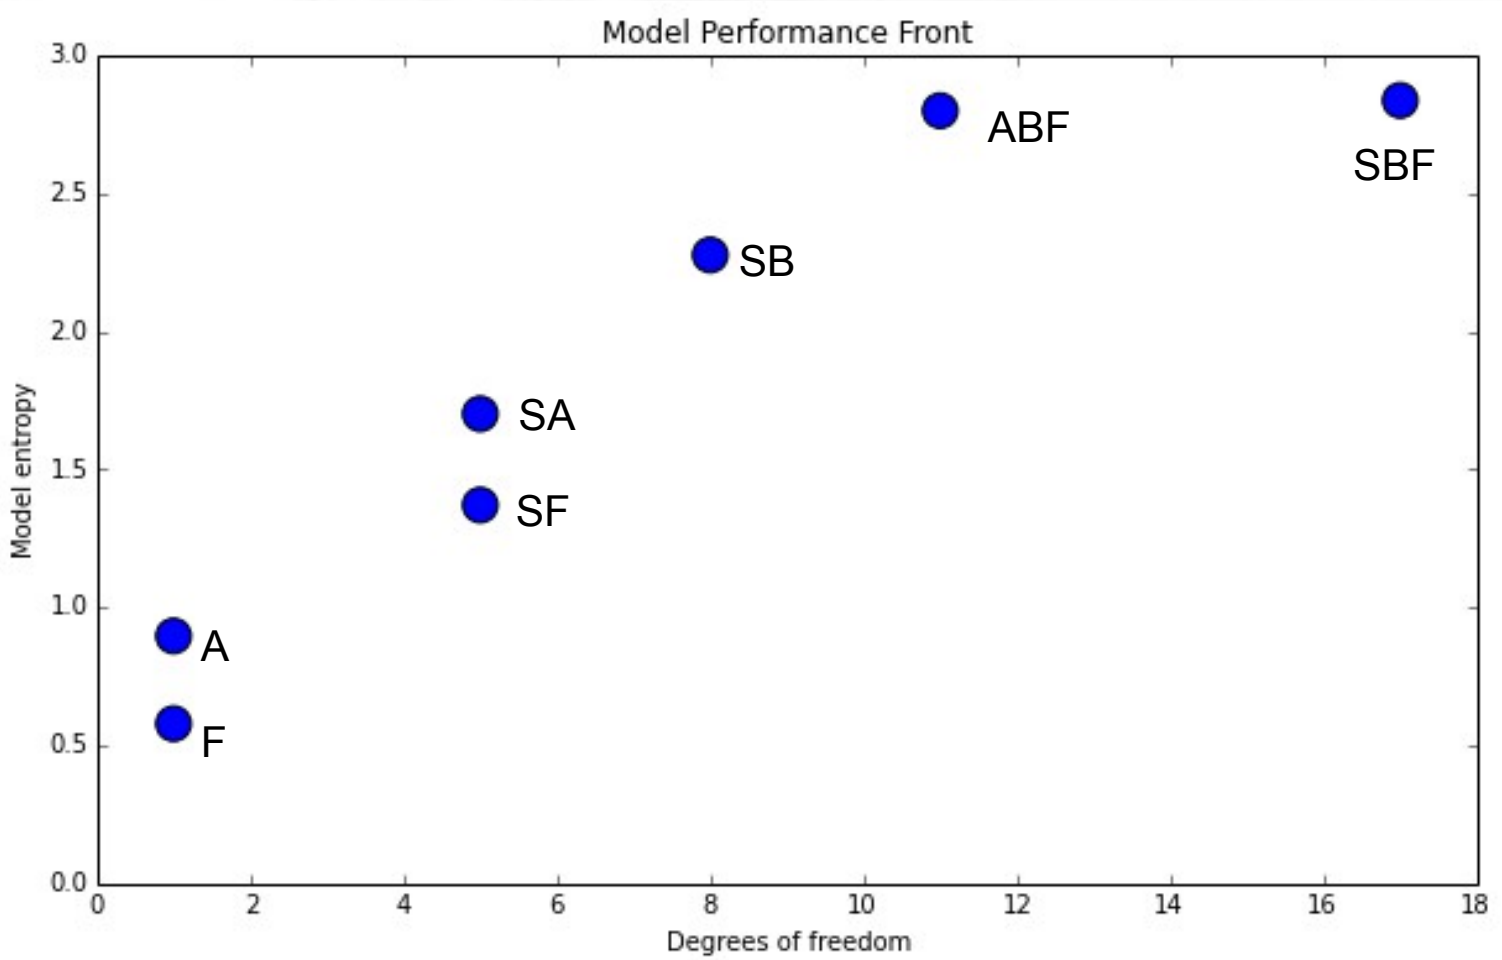

* Assumes all models are statistically significant.
* Sample size dimension not displayed, since all models have the same sample size

Predictions From a Model

- Given a positive, zero, or negative score, how long is your post?

Model “SB”

Score	B < 50	50<=B<130	B >= 130
S < 0	0.0045	0.0046	0.0026
S = 0	0.0490	0.0830	0.0682
S > 0	0.1156	0.2916	0.3810

Body Length

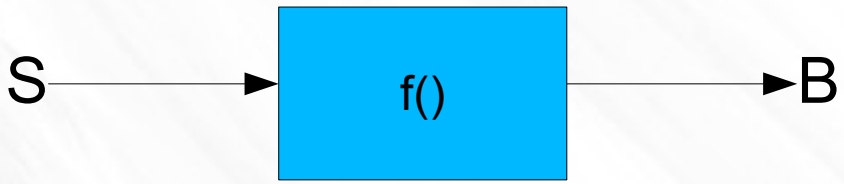

*Note issue of unbalanced distributions.

Predictions From a Model

- Given a zero score, how long is your post?

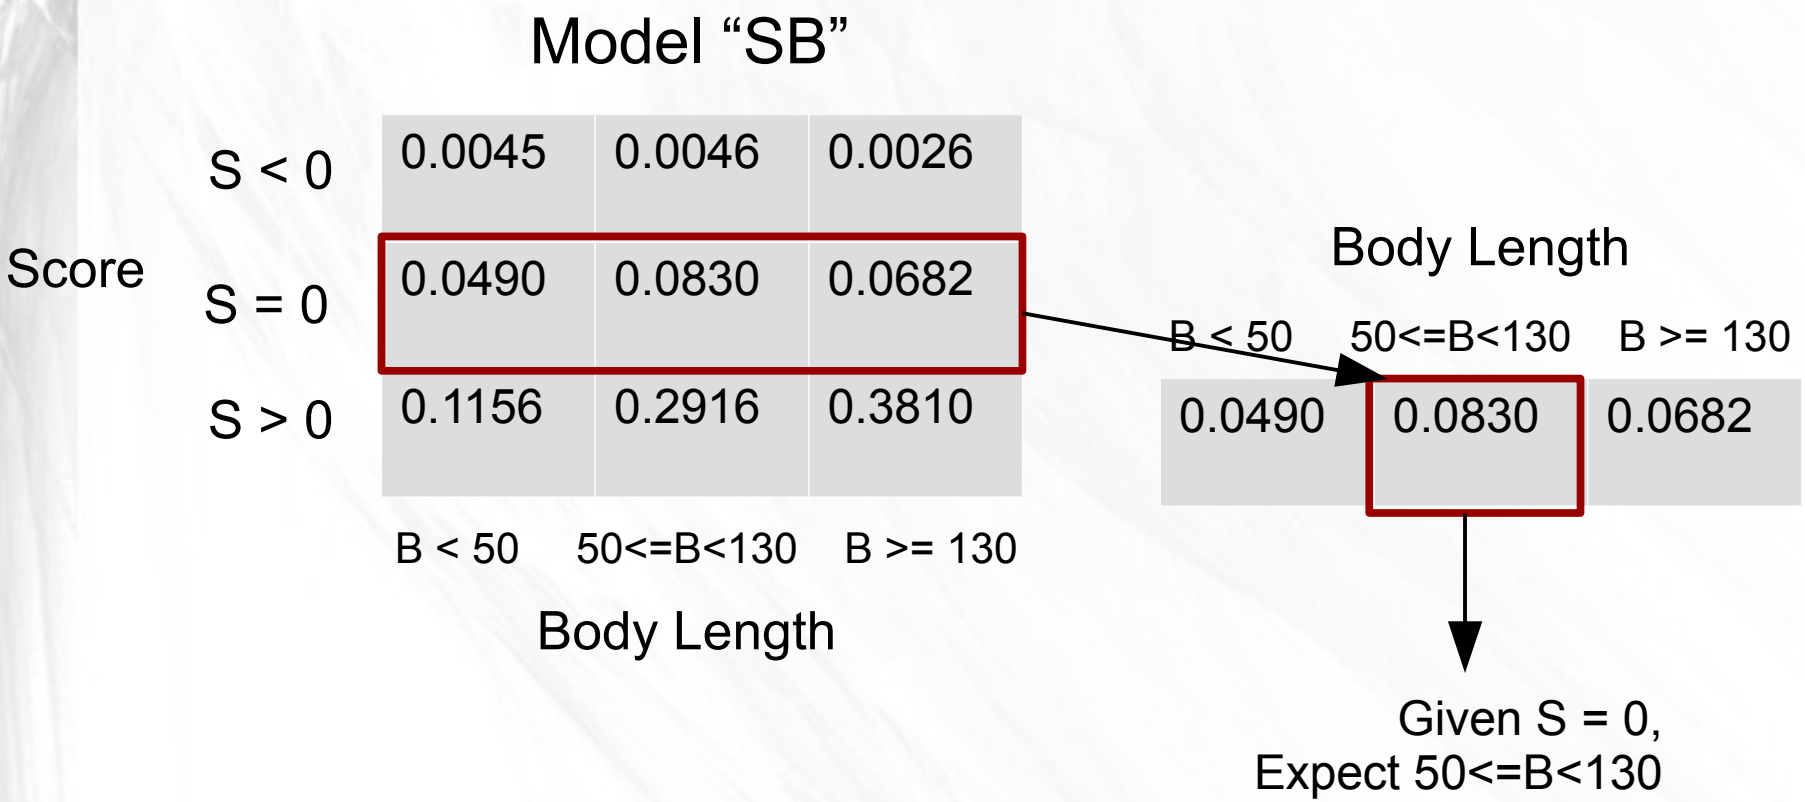

Predictions From a Model

- Given a positive score, how long is your post?

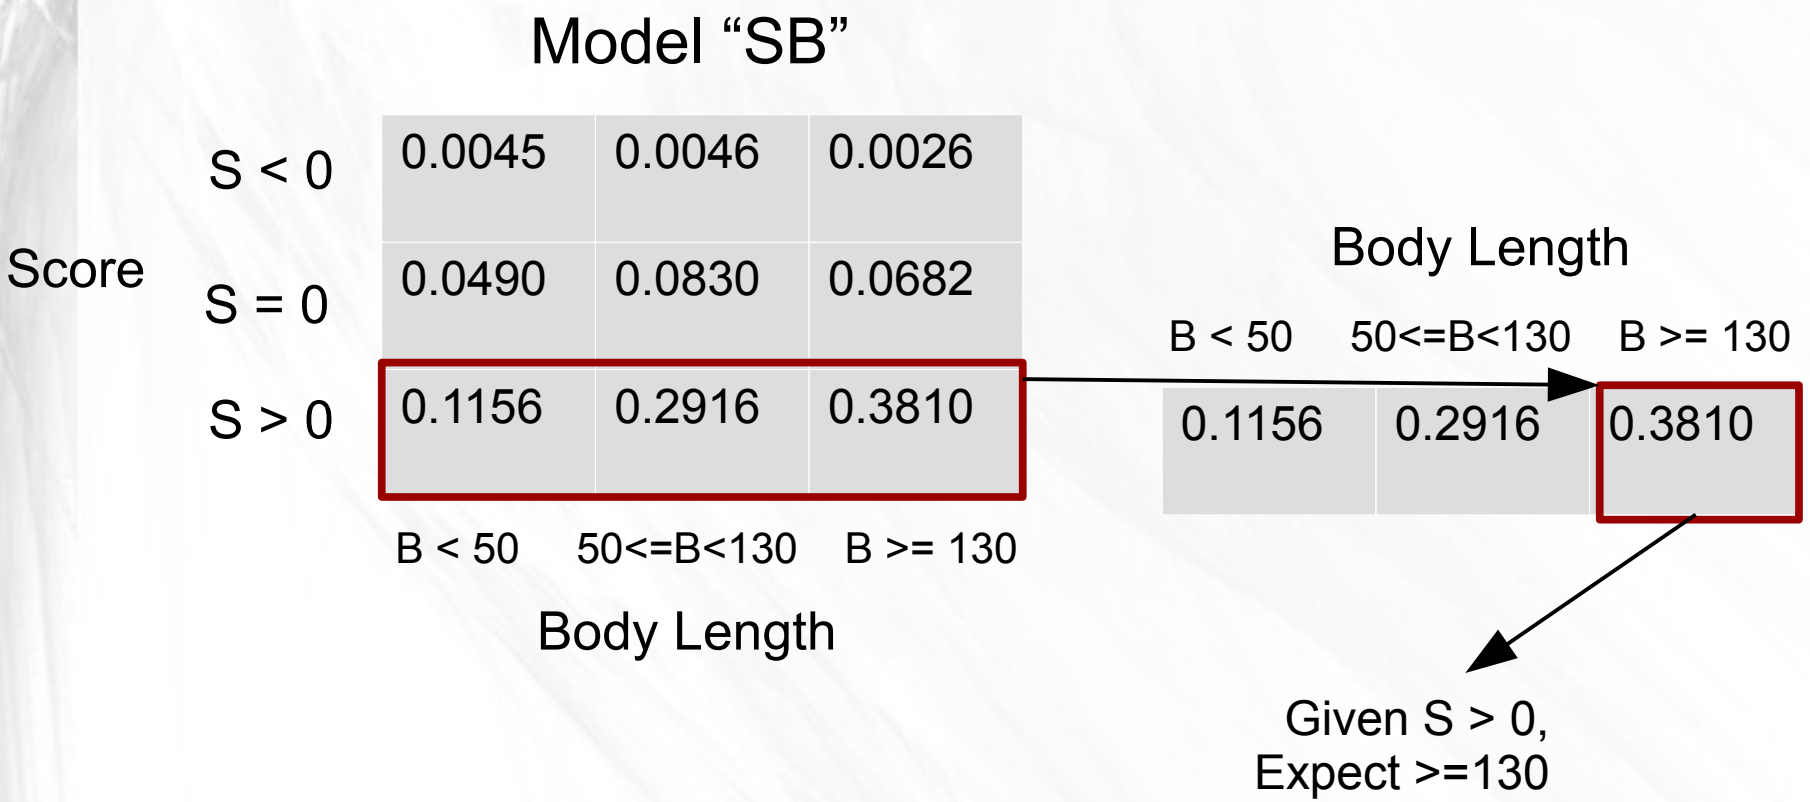

Credits, Other Resources

- www.github.com/TheGrimmScientist/DMM_Sim
  - Thanks to Ryan Price for working with me on that simulator and this example.
- http://occam.research.pdx.edu/occam/weboccam.cgi

Contact: TheGrimmScientist@gmail.com

Follow: @GrimmScientist